

# fer2013

April 27, 2023

```
[1]: ! pip install -q kaggle
from google.colab import files
files.upload()
! mkdir ~/.kaggle
! cp kaggle.json ~/.kaggle/
! chmod 600 ~/.kaggle/kaggle.json
```

<IPython.core.display.HTML object>

Saving kaggle.json to kaggle.json

```
[2]: !kaggle datasets download -d msambare/fer2013
```

Downloading fer2013.zip to /content  
76% 46.0M/60.3M [00:00<00:00, 94.9MB/s]  
100% 60.3M/60.3M [00:00<00:00, 114MB/s]

```
[ ]: !unzip fer2013.zip
```

```
[19]: import os
import numpy as np
import cv2 as cv
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.layers import Conv2D, MaxPool2D, Dropout, BatchNormalization, Dense,
↳ Flatten, GlobalAveragePooling2D
from keras.models import Sequential
from keras.applications.vgg16 import VGG16
from keras.applications.efficientnet import EfficientNetB4
from tensorflow.keras.applications.resnet50 import ResNet50
import xgboost as xgb
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
↳ confusion_matrix
```

```
[5]: def input_data(folder_path, output_data):  
    for dirs in os.listdir(folder_path):  
        class_name = dirs  
        new_path = os.path.join(folder_path, class_name)  
        for img in os.listdir(new_path):  
            img_arr = cv.imread(os.path.join(new_path, img))  
            output_data.append([img_arr, class_name])  
    return output_data
```

```
[6]: train_data = input_data("/content/train", [])
```

```
[7]: test_data = input_data("/content/test", [])
```

```
[8]: np.random.shuffle(train_data)  
np.random.shuffle(test_data)
```

```
[9]: train_images = []  
train_labels = []  
for features, labels in train_data:  
    train_images.append(features)  
    train_labels.append(labels)
```

```
[10]: test_images = []  
test_labels = []  
for features, labels in test_data:  
    test_images.append(features)  
    test_labels.append(labels)
```

```
[11]: label_enc = LabelEncoder()  
train_labels = label_enc.fit_transform(train_labels)  
test_labels = label_enc.transform(test_labels)
```

```
[12]: train_images = np.array(train_images)  
train_labels = np.array(train_labels)  
test_images = np.array(test_images)  
test_labels = np.array(test_labels)
```

```
[13]: print(f"Shape of the train images {train_images.shape}")  
print(f"Shape of the train labels {train_labels.shape}")  
print(f"Shape of the test images {test_images.shape}")  
print(f"Shape of the test labels {test_labels.shape}")
```

```
Shape of the train images (28709, 48, 48, 3)  
Shape of the train labels (28709,)  
Shape of the test images (7178, 48, 48, 3)  
Shape of the test labels (7178,)
```

```
[14]: train_images = train_images/255
      test_images = test_images/255
```

```
[17]: plt.figure(figsize=(15,10))
      for i in range(25):
          plt.subplot(5, 5, i+1)
          plt.imshow(test_images[i])
          plt.title(f"{label_enc.inverse_transform([test_labels[i]])}")
          plt.axis("off")
```



## CNN Model

```
[26]: model1 = Sequential()
```

```
[27]: model1.add(Conv2D(32, (3, 3), input_shape=(48,48,3), activation="leaky_relu"))
      model1.add(MaxPool2D(2,2))
      model1.add(Conv2D(64, (3, 3), activation="leaky_relu"))
      model1.add(MaxPool2D(2,2))
      model1.add(Conv2D(128, (3, 3), activation="leaky_relu"))
      model1.add(MaxPool2D(2,2))
      model1.add(Conv2D(256, (3, 3), activation="leaky_relu"))
      model1.add(MaxPool2D(2,2))
```

```
model1.add(Flatten())
model1.add(Dense(512, activation="relu"))
model1.add(Dense(7, activation="softmax"))
```

```
[54]: model1.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 46, 46, 32)	896
max_pooling2d_8 (MaxPooling2D)	(None, 23, 23, 32)	0
conv2d_10 (Conv2D)	(None, 21, 21, 64)	18496
max_pooling2d_9 (MaxPooling2D)	(None, 10, 10, 64)	0
conv2d_11 (Conv2D)	(None, 8, 8, 128)	73856
max_pooling2d_10 (MaxPooling2D)	(None, 4, 4, 128)	0
conv2d_12 (Conv2D)	(None, 2, 2, 256)	295168
max_pooling2d_11 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten_2 (Flatten)	(None, 256)	0
dense_2 (Dense)	(None, 512)	131584
dense_3 (Dense)	(None, 7)	3591
Total params: 523,591		
Trainable params: 523,591		
Non-trainable params: 0		

```
[28]: model1.compile(optimizer="adam", loss="sparse_categorical_crossentropy",
    ↪ metrics = ['accuracy'])
```

```
[29]: history1 = model1.fit(train_images, train_labels, validation_data=(test_images,
    ↪ test_labels), epochs=30)
```

Epoch 1/30  
898/898 [=====] - 9s 8ms/step - loss: 1.6083 -  
accuracy: 0.3585 - val\_loss: 1.3817 - val\_accuracy: 0.4709

Epoch 2/30  
898/898 [=====] - 5s 6ms/step - loss: 1.2594 -  
accuracy: 0.5186 - val\_loss: 1.2106 - val\_accuracy: 0.5404

Epoch 3/30  
898/898 [=====] - 6s 6ms/step - loss: 1.0990 -  
accuracy: 0.5833 - val\_loss: 1.1550 - val\_accuracy: 0.5626

Epoch 4/30  
898/898 [=====] - 5s 6ms/step - loss: 0.9530 -  
accuracy: 0.6392 - val\_loss: 1.1733 - val\_accuracy: 0.5639

Epoch 5/30  
898/898 [=====] - 6s 6ms/step - loss: 0.7959 -  
accuracy: 0.7029 - val\_loss: 1.2274 - val\_accuracy: 0.5766

Epoch 6/30  
898/898 [=====] - 5s 6ms/step - loss: 0.6272 -  
accuracy: 0.7686 - val\_loss: 1.2931 - val\_accuracy: 0.5812

Epoch 7/30  
898/898 [=====] - 6s 6ms/step - loss: 0.4681 -  
accuracy: 0.8288 - val\_loss: 1.4943 - val\_accuracy: 0.5724

Epoch 8/30  
898/898 [=====] - 5s 6ms/step - loss: 0.3515 -  
accuracy: 0.8708 - val\_loss: 1.8871 - val\_accuracy: 0.5659

Epoch 9/30  
898/898 [=====] - 6s 7ms/step - loss: 0.2775 -  
accuracy: 0.9010 - val\_loss: 2.0380 - val\_accuracy: 0.5607

Epoch 10/30  
898/898 [=====] - 5s 6ms/step - loss: 0.2231 -  
accuracy: 0.9215 - val\_loss: 2.3632 - val\_accuracy: 0.5702

Epoch 11/30  
898/898 [=====] - 5s 6ms/step - loss: 0.1971 -  
accuracy: 0.9298 - val\_loss: 2.5455 - val\_accuracy: 0.5652

Epoch 12/30  
898/898 [=====] - 6s 6ms/step - loss: 0.1841 -  
accuracy: 0.9368 - val\_loss: 2.7169 - val\_accuracy: 0.5681

Epoch 13/30  
898/898 [=====] - 5s 6ms/step - loss: 0.1562 -  
accuracy: 0.9466 - val\_loss: 2.8244 - val\_accuracy: 0.5571

Epoch 14/30  
898/898 [=====] - 6s 6ms/step - loss: 0.1648 -  
accuracy: 0.9453 - val\_loss: 2.8391 - val\_accuracy: 0.5634

Epoch 15/30  
898/898 [=====] - 5s 6ms/step - loss: 0.1431 -  
accuracy: 0.9527 - val\_loss: 2.8315 - val\_accuracy: 0.5571

Epoch 16/30  
898/898 [=====] - 6s 6ms/step - loss: 0.1375 -  
accuracy: 0.9539 - val\_loss: 3.0607 - val\_accuracy: 0.5617

```

Epoch 17/30
898/898 [=====] - 5s 6ms/step - loss: 0.1251 -
accuracy: 0.9598 - val_loss: 3.2240 - val_accuracy: 0.5666
Epoch 18/30
898/898 [=====] - 5s 6ms/step - loss: 0.1252 -
accuracy: 0.9590 - val_loss: 3.1433 - val_accuracy: 0.5723
Epoch 19/30
898/898 [=====] - 5s 6ms/step - loss: 0.1236 -
accuracy: 0.9600 - val_loss: 3.4812 - val_accuracy: 0.5736
Epoch 20/30
898/898 [=====] - 5s 6ms/step - loss: 0.1328 -
accuracy: 0.9571 - val_loss: 3.3832 - val_accuracy: 0.5729
Epoch 21/30
898/898 [=====] - 6s 6ms/step - loss: 0.1155 -
accuracy: 0.9627 - val_loss: 3.7177 - val_accuracy: 0.5585
Epoch 22/30
898/898 [=====] - 5s 6ms/step - loss: 0.1088 -
accuracy: 0.9640 - val_loss: 4.0101 - val_accuracy: 0.5651
Epoch 23/30
898/898 [=====] - 6s 6ms/step - loss: 0.1280 -
accuracy: 0.9592 - val_loss: 3.7505 - val_accuracy: 0.5665
Epoch 24/30
898/898 [=====] - 5s 6ms/step - loss: 0.1065 -
accuracy: 0.9661 - val_loss: 4.0427 - val_accuracy: 0.5705
Epoch 25/30
898/898 [=====] - 5s 6ms/step - loss: 0.1044 -
accuracy: 0.9679 - val_loss: 3.8090 - val_accuracy: 0.5649
Epoch 26/30
898/898 [=====] - 5s 6ms/step - loss: 0.1129 -
accuracy: 0.9628 - val_loss: 3.4479 - val_accuracy: 0.5617
Epoch 27/30
898/898 [=====] - 5s 6ms/step - loss: 0.1135 -
accuracy: 0.9647 - val_loss: 3.8520 - val_accuracy: 0.5663
Epoch 28/30
898/898 [=====] - 6s 6ms/step - loss: 0.0980 -
accuracy: 0.9698 - val_loss: 3.5648 - val_accuracy: 0.5634
Epoch 29/30
898/898 [=====] - 8s 8ms/step - loss: 0.1123 -
accuracy: 0.9656 - val_loss: 4.0841 - val_accuracy: 0.5667
Epoch 30/30
898/898 [=====] - 9s 10ms/step - loss: 0.1017 -
accuracy: 0.9682 - val_loss: 4.3113 - val_accuracy: 0.5715

```


```

[31]: plt.plot(history1.history["loss"])
      plt.plot(history1.history["val_loss"])
      plt.xlabel("epochs")
      plt.ylabel("loss")

```

```
plt.legend(["train","valid"])
```

```
[31]: <matplotlib.legend.Legend at 0x7fc4b88452e0>
```



```
[32]: y_pred1 = model1.predict(test_images)
```


```
225/225 [=====] - 1s 2ms/step
```

```
[33]: print(classification_report(y_pred1.argmax(axis=1), test_labels))
```

	precision	recall	f1-score	support
0	0.52	0.44	0.48	1139
1	0.49	0.64	0.55	85
2	0.41	0.42	0.41	1021
3	0.78	0.75	0.76	1836
4	0.50	0.53	0.52	1166
5	0.44	0.46	0.45	1181
6	0.70	0.78	0.74	750
accuracy			0.57	7178
macro avg	0.55	0.57	0.56	7178

weighted avg      0.57      0.57      0.57      7178

```
[37]: sns.heatmap(confusion_matrix(y_pred1.argmax(axis=1), test_labels), fmt='g',
    ↪annot=True)
plt.show()
```



CNN Model with Dropout, Global Average Pooling2D and BatchNormalization layer

```
[38]: model2 = Sequential()
```

```
[39]: model2.add(Conv2D(32, (3, 3), input_shape=(48,48,3), activation="leaky_relu"))
model2.add(BatchNormalization())
model2.add(Dropout(0.1))
model2.add(MaxPool2D(2,2))
model2.add(Conv2D(64, (3, 3), activation="leaky_relu"))
model2.add(BatchNormalization())
model2.add(Dropout(0.2))
model2.add(MaxPool2D(2,2))
model2.add(Conv2D(128, (3, 3), activation="leaky_relu"))
model2.add(BatchNormalization())
model2.add(Dropout(0.3))
```



```

model2.add(MaxPool2D(2,2))
model2.add(Conv2D(256, (3, 3), activation="leaky_relu"))
model2.add(MaxPool2D(2,2))
model2.add(BatchNormalization())
model2.add(GlobalAveragePooling2D())
model2.add(Dense(512, activation="relu"))
model2.add(BatchNormalization())
model2.add(Dense(7, activation="softmax"))

```

[55]: model2.summary()

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 46, 46, 32)	896
batch_normalization (Batch Normalization)	(None, 46, 46, 32)	128
dropout (Dropout)	(None, 46, 46, 32)	0
max_pooling2d_12 (MaxPooling2D)	(None, 23, 23, 32)	0
conv2d_14 (Conv2D)	(None, 21, 21, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 21, 21, 64)	256
dropout_1 (Dropout)	(None, 21, 21, 64)	0
max_pooling2d_13 (MaxPooling2D)	(None, 10, 10, 64)	0
conv2d_15 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 8, 8, 128)	512
dropout_2 (Dropout)	(None, 8, 8, 128)	0
max_pooling2d_14 (MaxPooling2D)	(None, 4, 4, 128)	0
conv2d_16 (Conv2D)	(None, 2, 2, 256)	295168

max_pooling2d_15 (MaxPoolin g2D)	(None, 1, 1, 256)	0
batch_normalization_3 (Batc hNormalization)	(None, 1, 1, 256)	1024
global_average_pooling2d (G lobalAveragePooling2D)	(None, 256)	0
dense_4 (Dense)	(None, 512)	131584
batch_normalization_4 (Batc hNormalization)	(None, 512)	2048
dense_5 (Dense)	(None, 7)	3591

```
=====
Total params: 527,559
Trainable params: 525,575
Non-trainable params: 1,984
-----
```

```
[40]: model2.compile(optimizer='adam', loss="sparse_categorical_crossentropy",  
    ↪metrics=['accuracy'])
```

```
[41]: history2 = model2.fit(train_images, train_labels, validation_data=(test_images,  
    ↪test_labels), epochs=30)
```


```
Epoch 1/30
898/898 [=====] - 14s 10ms/step - loss: 1.6135 -  
accuracy: 0.3894 - val_loss: 1.5822 - val_accuracy: 0.4030
Epoch 2/30
898/898 [=====] - 9s 10ms/step - loss: 1.3493 -  
accuracy: 0.4911 - val_loss: 1.3644 - val_accuracy: 0.4674
Epoch 3/30
898/898 [=====] - 9s 10ms/step - loss: 1.2384 -  
accuracy: 0.5365 - val_loss: 1.4760 - val_accuracy: 0.4695
Epoch 4/30
898/898 [=====] - 8s 9ms/step - loss: 1.1588 -  
accuracy: 0.5645 - val_loss: 1.3102 - val_accuracy: 0.5163
Epoch 5/30
898/898 [=====] - 9s 11ms/step - loss: 1.0906 -  
accuracy: 0.5911 - val_loss: 1.4555 - val_accuracy: 0.4503
Epoch 6/30
898/898 [=====] - 9s 10ms/step - loss: 1.0224 -  
accuracy: 0.6156 - val_loss: 1.3114 - val_accuracy: 0.5109
Epoch 7/30
898/898 [=====] - 9s 9ms/step - loss: 0.9615 -
```

accuracy: 0.6385 - val\_loss: 1.1964 - val\_accuracy: 0.5620  
 Epoch 8/30  
 898/898 [=====] - 8s 9ms/step - loss: 0.8951 -  
 accuracy: 0.6659 - val\_loss: 1.2569 - val\_accuracy: 0.5444  
 Epoch 9/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.8359 -  
 accuracy: 0.6886 - val\_loss: 1.4681 - val\_accuracy: 0.4941  
 Epoch 10/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.7763 -  
 accuracy: 0.7098 - val\_loss: 1.2236 - val\_accuracy: 0.5860  
 Epoch 11/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.7243 -  
 accuracy: 0.7297 - val\_loss: 1.3832 - val\_accuracy: 0.5627  
 Epoch 12/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.6722 -  
 accuracy: 0.7470 - val\_loss: 1.2763 - val\_accuracy: 0.5872  
 Epoch 13/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.6255 -  
 accuracy: 0.7688 - val\_loss: 1.4210 - val\_accuracy: 0.5644  
 Epoch 14/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.5943 -  
 accuracy: 0.7818 - val\_loss: 1.3956 - val\_accuracy: 0.5756  
 Epoch 15/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.5550 -  
 accuracy: 0.7958 - val\_loss: 1.4557 - val\_accuracy: 0.5780  
 Epoch 16/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.5178 -  
 accuracy: 0.8093 - val\_loss: 1.5094 - val\_accuracy: 0.5502  
 Epoch 17/30  
 898/898 [=====] - 10s 11ms/step - loss: 0.4985 -  
 accuracy: 0.8155 - val\_loss: 1.4877 - val\_accuracy: 0.5701  
 Epoch 18/30  
 898/898 [=====] - 10s 11ms/step - loss: 0.4597 -  
 accuracy: 0.8312 - val\_loss: 1.4873 - val\_accuracy: 0.5802  
 Epoch 19/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.4444 -  
 accuracy: 0.8374 - val\_loss: 1.5455 - val\_accuracy: 0.5885  
 Epoch 20/30  
 898/898 [=====] - 10s 11ms/step - loss: 0.4346 -  
 accuracy: 0.8402 - val\_loss: 1.6302 - val\_accuracy: 0.5868  
 Epoch 21/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.4050 -  
 accuracy: 0.8508 - val\_loss: 1.6066 - val\_accuracy: 0.5846  
 Epoch 22/30  
 898/898 [=====] - 9s 10ms/step - loss: 0.3907 -  
 accuracy: 0.8591 - val\_loss: 1.5995 - val\_accuracy: 0.5933  
 Epoch 23/30  
 898/898 [=====] - 8s 9ms/step - loss: 0.3766 -

```
accuracy: 0.8634 - val_loss: 1.6534 - val_accuracy: 0.5743
Epoch 24/30
898/898 [=====] - 9s 10ms/step - loss: 0.3686 -
accuracy: 0.8647 - val_loss: 1.6265 - val_accuracy: 0.5851
Epoch 25/30
898/898 [=====] - 9s 10ms/step - loss: 0.3427 -
accuracy: 0.8740 - val_loss: 1.8270 - val_accuracy: 0.5871
Epoch 26/30
898/898 [=====] - 9s 10ms/step - loss: 0.3363 -
accuracy: 0.8770 - val_loss: 1.7140 - val_accuracy: 0.5907
Epoch 27/30
898/898 [=====] - 10s 11ms/step - loss: 0.3417 -
accuracy: 0.8749 - val_loss: 1.7387 - val_accuracy: 0.5981
Epoch 28/30
898/898 [=====] - 10s 11ms/step - loss: 0.3181 -
accuracy: 0.8845 - val_loss: 1.9459 - val_accuracy: 0.5690
Epoch 29/30
898/898 [=====] - 10s 11ms/step - loss: 0.3087 -
accuracy: 0.8898 - val_loss: 1.7662 - val_accuracy: 0.5855
Epoch 30/30
898/898 [=====] - 14s 15ms/step - loss: 0.3047 -
accuracy: 0.8900 - val_loss: 1.7754 - val_accuracy: 0.5896
```

```
[42]: plt.plot(history2.history["loss"])
plt.plot(history2.history["val_loss"])
plt.xlabel("epochs")
plt.ylabel("loss")
plt.legend(["train", "valid"])
```

```
[42]: <matplotlib.legend.Legend at 0x7fc424199e50>
```




```
[43]: y_pred2 = model2.predict(test_images)
```

```
225/225 [=====] - 1s 4ms/step
```

```
[44]: print(classification_report(y_pred2.argmax(axis=1), test_labels))
```

	precision	recall	f1-score	support
0	0.50	0.52	0.51	920
1	0.57	0.56	0.56	113
2	0.46	0.41	0.44	1150
3	0.73	0.86	0.79	1507
4	0.55	0.54	0.54	1249
5	0.53	0.44	0.48	1521
6	0.70	0.81	0.75	718
accuracy			0.59	7178
macro avg	0.58	0.59	0.58	7178
weighted avg	0.58	0.59	0.58	7178

```
[45]: sns.heatmap(confusion_matrix(y_pred2.argmax(axis=1), test_labels), fmt='g',
    ↪annot=True)
plt.show()
```



### CNN Model with Dropout and GlobalAveragePooling2D layer

```
[62]: model3 = Sequential()
```

```
[63]: model3.add(Conv2D(32, (3, 3), input_shape=(48,48,3), activation="leaky_relu",
    ↪padding="same"))
model3.add(Conv2D(32, (3, 3), activation="leaky_relu", padding="same"))
model3.add(Dropout(0.1))
model3.add(MaxPool2D(2,2))
model3.add(Conv2D(64, (3, 3), activation="leaky_relu", padding="same"))
model3.add(Dropout(0.2))
model3.add(MaxPool2D(2,2))
model3.add(Conv2D(128, (3, 3), activation="leaky_relu", padding="same"))
model3.add(Dropout(0.3))
model3.add(MaxPool2D(2,2))
model3.add(Conv2D(256, (3, 3), activation="leaky_relu", padding="same"))
model3.add(MaxPool2D(2,2))
model3.add(Dropout(0.3))
```

```

model3.add(GlobalAveragePooling2D())
model3.add(Dense(512, activation="relu"))
model3.add(Dense(7, activation="softmax"))

```

```
[64]: model3.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
conv2d_42 (Conv2D)	(None, 48, 48, 32)	896
conv2d_43 (Conv2D)	(None, 48, 48, 32)	9248
dropout_17 (Dropout)	(None, 48, 48, 32)	0
max_pooling2d_32 (MaxPooling2D)	(None, 24, 24, 32)	0
conv2d_44 (Conv2D)	(None, 24, 24, 64)	18496
dropout_18 (Dropout)	(None, 24, 24, 64)	0
max_pooling2d_33 (MaxPooling2D)	(None, 12, 12, 64)	0
conv2d_45 (Conv2D)	(None, 12, 12, 128)	73856
dropout_19 (Dropout)	(None, 12, 12, 128)	0
max_pooling2d_34 (MaxPooling2D)	(None, 6, 6, 128)	0
conv2d_46 (Conv2D)	(None, 6, 6, 256)	295168
max_pooling2d_35 (MaxPooling2D)	(None, 3, 3, 256)	0
dropout_20 (Dropout)	(None, 3, 3, 256)	0
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 256)	0
dense_10 (Dense)	(None, 512)	131584
dense_11 (Dense)	(None, 7)	3591

```
=====
Total params: 532,839
Trainable params: 532,839
Non-trainable params: 0
-----
```

```
[65]: model3.compile(optimizer='adam', loss="sparse_categorical_crossentropy",  
    ↪metrics=['accuracy'])
```

```
[66]: history3 = model3.fit(train_images, train_labels, validation_data =  
    ↪(test_images, test_labels), epochs=30)
```

```
Epoch 1/30
898/898 [=====] - 14s 12ms/step - loss: 1.7980 -  
accuracy: 0.2588 - val_loss: 1.7406 - val_accuracy: 0.2856
Epoch 2/30
898/898 [=====] - 9s 10ms/step - loss: 1.5542 -  
accuracy: 0.3918 - val_loss: 1.4392 - val_accuracy: 0.4570
Epoch 3/30
898/898 [=====] - 10s 11ms/step - loss: 1.3331 -  
accuracy: 0.4868 - val_loss: 1.3010 - val_accuracy: 0.5001
Epoch 4/30
898/898 [=====] - 9s 10ms/step - loss: 1.2165 -  
accuracy: 0.5331 - val_loss: 1.2635 - val_accuracy: 0.5293
Epoch 5/30
898/898 [=====] - 10s 11ms/step - loss: 1.1348 -  
accuracy: 0.5624 - val_loss: 1.2208 - val_accuracy: 0.5412
Epoch 6/30
898/898 [=====] - 9s 10ms/step - loss: 1.0856 -  
accuracy: 0.5869 - val_loss: 1.3450 - val_accuracy: 0.5287
Epoch 7/30
898/898 [=====] - 10s 11ms/step - loss: 1.0335 -  
accuracy: 0.6034 - val_loss: 1.1944 - val_accuracy: 0.5683
Epoch 8/30
898/898 [=====] - 10s 11ms/step - loss: 0.9938 -  
accuracy: 0.6234 - val_loss: 1.1370 - val_accuracy: 0.5800
Epoch 9/30
898/898 [=====] - 10s 11ms/step - loss: 0.9577 -  
accuracy: 0.6353 - val_loss: 1.2025 - val_accuracy: 0.5665
Epoch 10/30
898/898 [=====] - 10s 11ms/step - loss: 0.9276 -  
accuracy: 0.6456 - val_loss: 1.1428 - val_accuracy: 0.5853
Epoch 11/30
898/898 [=====] - 10s 11ms/step - loss: 0.8878 -  
accuracy: 0.6641 - val_loss: 1.2033 - val_accuracy: 0.5791
Epoch 12/30
898/898 [=====] - 10s 11ms/step - loss: 0.8676 -  
accuracy: 0.6724 - val_loss: 1.1935 - val_accuracy: 0.5972
```




Epoch 13/30  
898/898 [=====] - 9s 10ms/step - loss: 0.8433 -  
accuracy: 0.6815 - val\_loss: 1.1968 - val\_accuracy: 0.5940  
Epoch 14/30  
898/898 [=====] - 10s 11ms/step - loss: 0.8120 -  
accuracy: 0.6932 - val\_loss: 1.1252 - val\_accuracy: 0.6055  
Epoch 15/30  
898/898 [=====] - 9s 11ms/step - loss: 0.7882 -  
accuracy: 0.7036 - val\_loss: 1.2851 - val\_accuracy: 0.5750  
Epoch 16/30  
898/898 [=====] - 9s 11ms/step - loss: 0.7660 -  
accuracy: 0.7114 - val\_loss: 1.1333 - val\_accuracy: 0.6080  
Epoch 17/30  
898/898 [=====] - 9s 10ms/step - loss: 0.7499 -  
accuracy: 0.7165 - val\_loss: 1.2586 - val\_accuracy: 0.5953  
Epoch 18/30  
898/898 [=====] - 10s 11ms/step - loss: 0.7342 -  
accuracy: 0.7218 - val\_loss: 1.2454 - val\_accuracy: 0.6110  
Epoch 19/30  
898/898 [=====] - 10s 11ms/step - loss: 0.7198 -  
accuracy: 0.7275 - val\_loss: 1.3128 - val\_accuracy: 0.5926  
Epoch 20/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6998 -  
accuracy: 0.7351 - val\_loss: 1.2665 - val\_accuracy: 0.6037  
Epoch 21/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6881 -  
accuracy: 0.7390 - val\_loss: 1.2841 - val\_accuracy: 0.5993  
Epoch 22/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6683 -  
accuracy: 0.7483 - val\_loss: 1.3442 - val\_accuracy: 0.6142  
Epoch 23/30  
898/898 [=====] - 9s 10ms/step - loss: 0.6638 -  
accuracy: 0.7482 - val\_loss: 1.4800 - val\_accuracy: 0.5939  
Epoch 24/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6528 -  
accuracy: 0.7557 - val\_loss: 1.2864 - val\_accuracy: 0.6162  
Epoch 25/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6314 -  
accuracy: 0.7619 - val\_loss: 1.6107 - val\_accuracy: 0.5717  
Epoch 26/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6252 -  
accuracy: 0.7646 - val\_loss: 1.4124 - val\_accuracy: 0.5903  
Epoch 27/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6186 -  
accuracy: 0.7683 - val\_loss: 1.3777 - val\_accuracy: 0.6103  
Epoch 28/30  
898/898 [=====] - 10s 11ms/step - loss: 0.6033 -  
accuracy: 0.7747 - val\_loss: 1.5117 - val\_accuracy: 0.6163

```
Epoch 29/30
898/898 [=====] - 10s 11ms/step - loss: 0.5950 -
accuracy: 0.7748 - val_loss: 1.3413 - val_accuracy: 0.6154
Epoch 30/30
898/898 [=====] - 10s 11ms/step - loss: 0.5897 -
accuracy: 0.7807 - val_loss: 1.4010 - val_accuracy: 0.6106
```

```
[68]: plt.plot(history2.history["loss"])
plt.plot(history2.history["val_loss"])
plt.xlabel("epochs")
plt.ylabel("loss")
plt.legend(["train", "valid"])
```

```
[68]: <matplotlib.legend.Legend at 0x7fc4215bc760>
```



```
[67]: y_pred3 = model3.predict(test_images)
```


```
225/225 [=====] - 1s 4ms/step
```

```
[69]: print(classification_report(y_pred3.argmax(axis=1), test_labels))
```

```
precision    recall  f1-score   support
```

0	0.55	0.48	0.52	1090
1	0.57	0.48	0.52	131
2	0.46	0.46	0.46	1019
3	0.81	0.82	0.81	1758
4	0.50	0.61	0.55	1005
5	0.51	0.47	0.49	1357
6	0.76	0.77	0.77	818
accuracy			0.61	7178
macro avg	0.59	0.59	0.59	7178
weighted avg	0.61	0.61	0.61	7178

```
[70]: sns.heatmap(confusion_matrix(y_pred3.argmax(axis=1), test_labels), fmt='g',
    ↪annot=True)
plt.show()
```



```
[ ]:
```